



# *Modeling Resource Management in Concurrent Computing Systems*

Isaac D. Scherson (aka The Schark)

Dept. of Computer Science (Systems)

School of Information and Computer Science

University of California, Irvine

Irvine, CA 92697-3425

isaac@ics.uci.edu - [www.ics.uci.edu/~isaac](http://www.ics.uci.edu/~isaac)

# Acknowledgement

This work was started in 1994-95 and the following were the pioneers to whom we owe the ideas presented here.


- Piotr Chrzastowski-Wachtel - Institute of Informatics, Warsaw University
- Dinesh Ramanathan - Somewhere in the Silicon Valley
- Raghu Subramanian - Somewhere in the Silicon Valley

This work was supported in part through various grants from NASA, the NSF and the AFOSR

# Warning !

This is a thought-provoking presentation. The ideas are very simple and are presented to motivate a little formal thinking about the scheduling problem in Concurrent Computing Systems.

# Concurrent Computing OS Model



General Purpose Parallel Operating System

# The Programming Model

- There seems to be a convergence on the data parallel programming in the form of HPF and Fortran 90.
- In a survey of 120 parallel algorithms from three ACM Symposia on Theory of Computing (STOC), all were found to be data parallel.
- This preponderance of the data parallel programming model is perhaps because it allows one to express a large degree of parallelism, while retaining the single-thread-of-control philosophy of sequential programming.

# The Machine Model

- A natural way of executing a data parallel program is on a SIMD machine.
- BUT . . . data parallelism does not equal SIMD.
- A data parallel program may be executed on an asynchronous MIMD machine:
  - Chief advantage: It is possible to run several jobs on a MIMD machine simultaneously.
  - A MIMD machine does not force unnecessary synchronization after every instruction, or unnecessary sequentialization of non-interfering branches, as a SIMD machine does.

# SPMD Programs

- The combination of the data parallel programming model and a MIMD machine model is called a SPMD execution model
- SPMD stands for Single Program Multiple Data
  - All processors execute the same program, but may be at different stages of the program at a given time
- For the remainder of this presentation, our discussion is predicated on the SPMD execution model.

# Virtual Processors

- A program is expressed in a language that describes a virtual machine.
- For a data parallel program, the virtual machine consists of a (typically large) number of identical *virtual processors* (VPs), communicating through an interconnection network.
- For instance, the standard data parallel program to multiply two  $N \times N$  matrices may be viewed as a virtual machine consisting of  $N^2$  VPs communicating in a mesh.



# Virtual Processors (Cont'd)

- For many years now, the concept of virtual processors has been relegated to the status of a mere logical aid to programmers
- In our view, the notion of virtual processors should form the fundamental basis for the definition of a Concurrent Computer System and for the resource management strategies embedded in the OS.

# Redefining Massive Parallelism

- A Programmer sees only a VM with as many Virtual Processors as s/he needs.
- Solutions are cast in the space of the problem and not on that of the physical machine.
- The measure of Massive Parallelism is through the number of VPs the physical machine is capable of emulating per unit of time.
- HENCE . . . It does not matter what the physical machine looks like, provided it can emulate Concurrent Computing Virtual Machines.

# Resource Management Operations

- **Spatial scheduling:** A space sharing policy that defines which physical processor each VP is allocated to.
- **Temporal scheduling:** A time sharing policy that dictates how each physical processor switches between the execution of the VPs allocated to it.
- **Load balancing:** Static and/or Dynamic redistribution of VPs among physical processors responding to some predetermined objective function.

# Resource Management Operations

- **Memory and I/O problems:** Memory limitations and I/O bottlenecks may also be phrased in terms of VPs. Memory limitations occur when a processor does not have enough local memory to hold all the VPs allocated to it. I/O bottlenecks are most pronounced while loading (roll-in) the VPs of a job from the disk into the local memories of various processors.

# AND NOW ...

A game you can play in your free time .....

# Model of the System

- Let  $\Pi = \{\pi_1, \pi_2, \dots, \pi_n\}$ . denote the set of physical processors or processing elements (PEs).
- Interconnection network details are ignored, only assuming that it *guarantees* the completion of all communications within a reasonable time bound (not necessarily a constant, as in the PRAM model).
- Let  $\mathcal{J} = \{J_1, J_2, \dots, J_m\}$  denote the set of jobs in the system at some time.
- For notational convenience, job  $J_i$  is identified with its set of VPs  $\{P_{i1}, P_{i2}, \dots, P_{i|J_i|}\}$ .

# Handling Time

- Break time into **time slices**: discrete intervals of equal length.
- The time slice is assumed to be machine dependent and constant.
- Within each time slice, every  $\pi_i$  runs some VP chosen from the VPs allocated to it, or it idles.

# Handling Time (Cont'd)

- A PE  $\pi_i$  runs a VP of a Job until:
  - The VP issues a communication request which cannot be completed (data to be read is not ready), or
  - The time slice expires.
- The communication instruction is considered an indivisible instruction and time slices are extended accordingly.



# Execution of Programs

- *For simplicity of analysis*, it is assumed that jobs run indefinitely.
  - The OS schedules an arriving Job as if the job is going to last forever in the system
- This approximation reflects that the time slice is very small (in the order of milliseconds) compared to the execution times of the jobs (in the order of minutes).
- The assumption of infinite running time might be removed with negligible change to the final results, but with significant complications in the definitions and proofs.

# The 2 Types of Scheduling

- Given a job,  $J$ , as a set of  $|J|$  VPs, how is this job allocated over the  $n$  physical PEs of the concurrent computer so that the resources are efficiently utilized?
- More generally, how are VPs of the same and different jobs, say  $J_1, J_2, \dots, J_m$ , allocated over the PEs to efficiently utilize all system resources?
  - This is called the problem of spatial scheduling or spatial allocation.
- Intuitively. Spatial schedule: “Where to run the VPs of job  $J$ ?”. Temporal schedule: “When to run the VPs of job  $J$ ?”.

# ... and more formally

**Definition 1** A **spatial schedule** is an  $n$ -tuple allocation function (The terms **allocation** and **spatial scheduling** will be used interchangeably)

$$\mathcal{A} : \mathcal{J} \longrightarrow \mathbb{N}^{|\Pi|}$$

such that for any job  $J$ ,  $\mathcal{A}(J) = (a_1, a_2, \dots, a_n)$ ,  $a_i$  VPs of  $J$  are allocated to processor  $\pi_i \in \Pi$ ,  $(\sum_{i=1}^n a_i = |J|)$ .

# ... and more formally (Cont'd)

**Definition 2** A temporal schedule for a time slice  $t$  is defined as an  $n$ -tuple function,  $S(t)$ , such that  $S(t)[i]$  is the job run by  $\pi_i$  in the time slice  $t$ . A **temporal schedule**,  $\mathcal{S}$ , is a time ordered sequence of  $S(t)$  for all  $t > 0$ .

**Definition 3** Define a schedule,  $\mathcal{S}$ , to be **impartial** if every job  $J \in \mathcal{J}$  occurs infinitely many times in  $\mathcal{S}$ .

An impartial schedule ensures that every job in it completes.

# Example

Consider a machine with 5 PEs and 5 jobs:

$J_1$  requires 1 VP,

$J_2$  requires 4 VPs,

$J_3$  requires 1 VP,

$J_4$  requires 5 VPs,

$J_5$  requires 1 VP.

... suppose we allocate the VPs of the Jobs as follows ...

# Example (Graphical)

|         |         |         |         |         |
|---------|---------|---------|---------|---------|
| 4       |         |         | 4       |         |
| 3       | 4       | 5       | 4       | 2       |
| 2       | 2       | 2       | 4       | 1       |
| $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |

With the following allocation functions:

$$\mathcal{A}(J_1) = (0, 0, 0, 0, 1)$$

$$\mathcal{A}(J_2) = (1, 1, 1, 0, 1)$$

$$\mathcal{A}(J_3) = (1, 0, 0, 0, 0)$$

$$\mathcal{A}(J_4) = (1, 1, 0, 3, 0)$$

$$\mathcal{A}(J_5) = (0, 0, 1, 0, 0)$$

# Job Execution

- **Global communication:** all VPs of a given Job execute a communication instruction.
- **Legal execution of a Job:** no VP of the Job is ahead of any other VP of the same job by more than one global communication step at any time.
- **Legal scheduling strategy:** will always choose any trailing VP before any non-trailing VP.
- A **step** begins when all VPs of a job have executed the same number of global communications and it ends at the occurrence of the same condition at a future time slice (after at least one VP of the job has been executed).

# A Job's Schedule Vector

**Definition 4** Define a job  $J$ 's  $t$ -schedule vector,  $\epsilon(t, J)$  as a 0-1  $n$ -tuple such that,

$$\epsilon(t, J)[k] = \begin{cases} 1 & \text{if } \pi_k \in \Pi \text{ runs job } J \text{ in time slice } t \\ 0 & \text{otherwise} \end{cases}$$



# A Job's Progress Vector

**Definition 5** Define progress of a job (at the beginning of a time slice  $t$ ) as an  $n$ -tuple function

$$\mathcal{P} : t \times \mathcal{J} \longrightarrow \mathbb{N}^{\Pi}$$

such that:

1.  $\mathcal{P}(t, J) = \vec{0}$  for  $t = 0$ ,
2. If  $\mathcal{P}(t - 1, J) + \epsilon(t, J) = \mathcal{A}(J)$  then  $\mathcal{P}(t, J) = \vec{0}$ ,  
else  $\mathcal{P}(t, J) = \mathcal{P}(t - 1, J) + \epsilon(t, J)$  (where the addition is componentwise)
3.  $\mathcal{P}(t, J) \leq \mathcal{A}(J)$  where the  $\leq$  order is applied componentwise.

# Progress Vector in English

The intuition behind the progress vector is as follows:

Start at the first step of a job  $J$ , let job  $J$  proceed on the processor's having 1's on the  $\epsilon(t, J)$ . If the progress vector reaches  $\mathcal{A}(J)$  then another step of job  $J$  is complete and progress counters associated with each processor are reset. A new step cannot be started until the previous one is complete. Every job has a progress vector associated with it at the beginning of every time slice.

# A more complete Example

Consider an allocation on 5 processors (same example):

|         |         |         |         |         |
|---------|---------|---------|---------|---------|
| 4       |         |         | 4       |         |
| 3       | 4       | 5       | 4       | 2       |
| 2       | 2       | 2       | 4       | 1       |
| $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |

And consider the problem of scheduling  $J_4$ .

# A more complete Example (Cont'd)

An arbitrary schedule for job  $J_4$  is as follows:

$$\mathcal{A}(J_4) = (1, 1, 0, 3, 0)$$

$$P(0, J_4) = (0, 0, 0, 0, 0)$$

$$\epsilon(1, J_4) = (0, 1, 0, 1, 0)$$

$$\left[ \begin{array}{lcl} P(0, J_4) + \epsilon(1, J_4) & = & P(1, J_4) = (0, 1, 0, 1, 0) \\ & & \epsilon(2, J_4) = (1, 0, 0, 1, 0) \end{array} \right.$$

$\epsilon(1, J_4) = (0, 1, 0, 1, 0)$  indicates that processors  $\pi_2$  and  $\pi_4$  run job  $J_4$ .  $\epsilon(1, J_4)$  is added to  $P(0, J_4)$  componentwise to give the new progress vector  $P(1, J_4)$  for job  $J_4$  and so on.

# A more complete Example (Cont'd)

Continuing to schedule  $J_4 \dots$

$$\left[ \begin{array}{lcl} P(1, J_4) + \epsilon(2, J_4) & = & P(2, J_4) = (1, 1, 0, 2, 0) \\ & & \epsilon(3, J_4) = (0, 0, 0, 1, 0) \end{array} \right.$$

$$P(2, J_4) + \epsilon(3, J_4) = (1, 1, 0, 3, 0)$$

Now,  $P(3, J_4) = \mathcal{A}(J_4)$ , so reset,  $P(3, J_4) = (0, 0, 0, 0, 0)$

# System's Progress and Schedule

**Definition 6** Define the **system progress matrix**,  $Q$ , as a  $(n \times m)$ -array of progress vectors (columns) of all jobs (rows) in the system.

**Definition 7** Define the **system schedule matrix**,  $S$ , as a  $n$  — column matrix such that each column corresponds to a physical processor while each row corresponds to a time slice. Each entry is numbered with the job number whose VP is executed at that time slice.

# A Schedule for the Example

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 2       | 2       | 2       | 4       | 1       |
| 2       | 3       | 4       | 5       | 4       | 2       |
| 3       | 4       |         |         | 4       |         |

The same three lines can be repeated until all jobs terminate.

The above will be dubbed the **Scheduling Matrix**

**HINT:** We'll be seeing periodic schedules.

# Another Schedule for the Example

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 2       | 2       | 2       | 4       | 2       |
| 2       | 3       | 4       | 5       | 4       | 1       |
| 3       | 4       |         | 5       | 4       | 1       |
| 4       | 2       | 2       | 2       | 4       | 2       |
| 5       | 3       | 4       | 5       | 4       | 1       |
| 6       | 4       |         | 5       | 4       | 1       |

Same period as the previous schedule (3) but only one idle slice per period.



# Schedules and Allocations

- A Schedule will strongly depend upon the initial Allocation.
- Challenge: Find another allocation/schedule for the example such that there are no idle time slices.

# Metrics

There are two points of view when measuring the system's performance.

- One focuses on the functional quality from the system manager's point of view, who is concerned with the throughput and utilization of the machine.
- The second focuses on the functional quality from the users' point of view who expects the system to respond within a specific time.

# Scheduling Metrics

For the case of Scheduling:

- Question: Which schedule is the “best”?
- The parallel OS should be able to satisfy two objectives:
  - Minimize processor idling time (important for the System).
  - Guarantee an upper limit on the system response time to every job (important to the single user).

# Idling Ratio

**Definition 8** For a  $n$ -processor system, for a schedule  $\mathcal{S}$ , and over an interval of  $\Delta t$  time slices, define an **idling ratio** or **throughput**

$i_{\mathcal{S}} : \mathbb{N} \rightarrow \mathbb{R}$  as:

$$i_{\mathcal{S}}(\Delta t) = \frac{\sum_{k \in \Delta t} i_k}{n \times \Delta t}$$

where  $i_k$  is the total number of idling processors during time slice  $k$ .

The idling ratio is a measure of the resource utilization.

# Idling Ratio for the Example

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 2       | 2       | 2       | 4       | 1       |
| 2       | 3       | 4       | 5       | 4       | 2       |
| 3       | 4       |         |         | 4       |         |

Idling Ratio for this schedule over the period of 3 time slices =

$$3/15 = 1/5 = 20\%$$

# Idling Ratio for the Example

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 2       | 2       | 2       | 4       | 2       |
| 2       | 3       | 4       | 5       | 4       | 1       |
| 3       | 4       |         | 5       | 4       | 1       |

Idling Ratio for this schedule over the period of 3 time slices =  $1/15$

# An Interesting Case

Consider a system where only one job is running and utilizes all processors ( $J_1$  has 5 VPs):

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 1       | 1       | 1       | 1       | 1       |

# An Interesting Case

...SUDDENLY ...



# An Interesting Case

...SUDDENLY ...

Another job  $J_2$  arrives and requires 4 VPs

# An Interesting Case

In a hurry, the system allocates and schedules  $J_2$  in the apparently obvious manner:

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 1       | 1       | 1       | 1       | 1       |
| 2       |         | 2       | 2       | 2       | 2       |

With an Idling Ratio of  $1/10 = 10\%$

# An Interesting Case

BUT ... Consider the following 0% Idling Ratio Schedule:

| Time/PE | $\pi_1$ | $\pi_2$ | $\pi_3$ | $\pi_4$ | $\pi_5$ |
|---------|---------|---------|---------|---------|---------|
| 1       | 1       | 2       | 2       | 2       | 2       |
| 2       | 1       | 2       | 2       | 2       | 2       |
| 3       | 1       | 2       | 2       | 2       | 2       |
| 4       | 1       | 2       | 2       | 2       | 2       |
| 5       | 1       | 2       | 2       | 2       | 2       |

The response time of  $J_1$  was sacrificed in favor of full 100% Physical Processor utilization. AN UNHAPPY USER !!!

# Happiness Function

**Definition 9** Define the **happiness function**,  $\bar{h}_{\Delta t}(J)$ , of a job  $J$  for a time length  $\Delta t$  as,

$$\bar{h}_{\Delta t}(J) = \min_{t \geq 0} \left\{ \frac{\sum_{\tau=t}^{t+\Delta t-1} \sum_{\pi \in \Pi} \epsilon(\tau, J)}{|J| \cdot \Delta t} \right\}$$

The time interval  $\Delta t$  is the **Impatience Latency** of the user.

# Happiness Function

- Intuitively, happiness represents the machine power that a job is given during a time slice. The term in the braces represents the fraction of the full parallel speedup that was achieved by job  $J$  in time interval  $(t, t + \Delta t - 1)$ .
- $\Delta t$  represents the expected **impatience latency** of the user. It is the minimum time the user is required to wait before the system responds.

# Virtual Happiness

- The proposed parallel OS paradigm allows the user to program in a virtual parallel model with no limit on the number of VPs in a job.
- To get a handle on the maximum parallel speedup that can be achieved for a job given  $n$  physical processors, define **virtual happiness**,  
$$v(J) = \min\{1, \frac{n}{|J|}\}.$$
- Virtual happiness of a job  $J$  represents the happiness of the job if the entire machine is used by  $J$ .
- Hence,  $J$  cannot be happier than  $v(J)$ .

# Happiness is a Balancing Act

- A Job (user) will be happier if it acquires more parallel speedup over time.
- If a schedule is unable to provide the required happiness for a job, rescheduling with a different allocation may have to be considered.
- If the required happiness for a set of jobs cannot be achieved by any allocation, then some jobs must be queued for later allocation.
- A happy schedule is not necessarily optimal for throughput. It merely guarantees an upper limit on the response time.

# Periodic Temporal Schedules

**Definition 10** For an allocation  $\mathcal{A}$ ,

$\mathcal{S}_p = S(t), S(t+1), \dots, S(t+p-1)$  is called a **periodic schedule**, iff  $\exists p > 0$  such that  $S(t) = S(t+p)$  for all  $t > t_s$ , where  $t_s$  is called the startup time of the schedule.  $p$  is the period of the schedule.

Also for each job  $J$  in  $\mathcal{S}_p$ ,

$$\sum_{1 \leq t \leq p, 1 \leq i \leq |\Pi|} [S(t)[i] = J] = k \cdot |J|$$

must be true for some positive integer  $k$  to ensure that each job executes at least  $k$  steps during each period ( $k$  need not be the same for every job).



# How far did we go?

- This is *Work in Progress*. A lot to be done !
- We have proven a few theorems on Impartial Periodic Schedules.
  - Periodic Schedules are the BEST !!
  - We have conditions to minimize Idling Ratio and maximize Happiness.

# Collaborations Welcome

Anybody interested ? ... Let's work together !!!

THANKS !!